

Apr 09, 17 14:09

main.c

Page 1/3

```

/*
 * Purpose: Time square matrix multiplication using CPU, GPU, cblas ddot,
 *          cblas daxpy, cublas ddot, cublas daxpy.
 *
 * Author: Gurpal Singh
 * Date: April 6, 2017
 * ME 571 Project 2
 */

#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <sys/resource.h>
#include <time.h>
#include "timer.h"
#include "gsl_cblas.h"
#include <cublas.h>

#include "cpumatrixmultiply.h"
#include "cpuddot.h"
#include "cpudaxpy.h"
#include "gpumatrixmultiply.h"
#include "gpuddot.h"
#include "gpudaxpy.h"

int main(void){

    //Initializing matrix Dimensions and scanning for values from user
    int m, n, k;

    printf("This program performs square matrix multiplication where A is 'm x n' and B is 'n x k'.\n");

    //Setting Matrix Dimensions
    printf("Enter the integer value for n: ");
    scanf("%d", &n);
    m = n;
    k = n;

    //Dynamic Memory Allocation for Matrices (I used flat arrays)
    //Matrix A
    double *a;
    cudaMallocManaged( &a, m * n * sizeof(double));

    //Matrix B
    double *b;
    cudaMallocManaged( &b, n * k * sizeof(double));

    //Matrix C
    double *c;
    cudaMallocManaged( &c, m * k * sizeof(double));

    //Initializing nmax and i and j used for loops
    int nmax = 19;
    int i, j;

    time_t t;
    srand( (unsigned) t );

    //Setting random values in Matrix A
    for (i = 0; i < m; i++){
        for (j = 0; j < n; j++){
            a[i*n+j] = rand() % (nmax+1);
        }
    }

    //Setting random values in Matrix B

```

Sunday April 09, 2017

main.c

Apr 09, 17 14:09

main.c

Page 2/3

```

    for (i = 0; i < n; i++){
        for (j = 0; j < k; j++){
            b[i*k+j] = rand() % (nmax+1);
        }
    }

    //Transposing and Storing as Column Major arrays
    double *acol;
    cudaMallocManaged( &acol, m * n * sizeof(double));
    for(i = 0; i < n; i++){
        for(j = 0; j < m; j++){
            acol[i*m+j] = a[j*n+i];
        }
    }

    double *bcol;
    cudaMallocManaged( &bcol, k * n * sizeof(double));
    for(int i = 0; i < k; i++){
        for( j = 0; j < n; j++){
            bcol[i*n+j] = b[j*k+i];
        }
    }

    //Timing the CPU Matrix Multiplication Method
    StartTimer();
    CPU_Matrix_Multiply(m, n, k, a, b, c);
    double CPU_time = GetTimer();
    CPU_time = CPU_time*1000; //Converting to ms

    //Timing the CPU Matrix Multiplication using cblas ddot method
    StartTimer();
    CPU_ddot(m, n, k, a, bcol, c);
    double CPU_ddot_time = GetTimer();
    CPU_ddot_time = CPU_ddot_time*1000; //Converting to ms

    //Timing the CPU Matrix Multiplication using cblas daxpy
    StartTimer();
    CPU_daxpy(m, n, k, acol, bcol, c);
    double CPU_daxpy_time = GetTimer();
    CPU_daxpy_time = CPU_daxpy_time*1000;

    //Parallel GPU Code Block and Grid Dimensions
    dim3 block(16,16);
    dim3 grid( (n+15)/16, (n+15)/16 );

    //Timing the GPU Matrix Multiplication Kernel
    cudaEvent_t timeStart, timeStop;

    cudaEventCreate(&timeStart);
    cudaEventCreate(&timeStop);
    float elapsedTime; //Has to be type float units ms

    cudaEventRecord(timeStart, 0);
    GPU_Matrix_Multiply_Kernel<<<grid, block>>> (a, b, c, n);
    cudaDeviceSynchronize();

    cudaEventRecord(timeStop, 0);
    cudaEventSynchronize(timeStop);
    cudaEventElapsedTime(&elapsedTime, timeStart, timeStop);

    cudaEventDestroy(timeStart);
    cudaEventDestroy(timeStop);

    //Timing the GPU cublas DDOT Matrix Multiplication Method

```

1/2

Apr 09, 17 14:09

main.c

Page 3/3

```

StartTimer();
GPU_ddot(m, n, k, a, bcol, c);
double GPU_ddot_time = GetTimer();
GPU_ddot_time = GPU_ddot_time*1000;

//Timing the GPU cublas DAXPY Matrix Multiplication Method
StartTimer();
GPU_daxpy(m, n, k, acol, bcol, c);
double GPU_daxpy_time = GetTimer();
GPU_daxpy_time = GPU_daxpy_time*1000;

//Writing The results to a file
FILE *fptr = fopen("Matrix_Multiply_Results.txt", "a+");

if (fptr == NULL) {
    printf("Error!");
    exit(1);
}

fprintf(fptr, "\n");
fprintf(fptr, "Matrix Size: %d\n", n);
fprintf(fptr, "elapsed wall time CPU Matrix Multiplication = %.3f ms\n", CPU_time);
fprintf(fptr, "elapsed wall time CPU cblas DDOT %.3f ms\n", CPU_ddot_time);

fprintf(fptr, "elapsed wall time CPU cblas DAXPY %.3f ms\n", CPU_daxpy_time);
fprintf(fptr, "elapsed wall time GPU Matrix Multiplication %.3f ms\n", elapsedTime);

fprintf(fptr, "elapsed wall time GPU cublas DDOT %.3f ms\n", GPU_ddot_time);
fprintf(fptr, "elapsed wall time GPU cublas DAXPY %.3f ms\n", GPU_daxpy_time);
fclose(fptr);

//Clean Up
cudaFree(a);
cudaFree(b);
cudaFree(c);
cudaFree(acol);
cudaFree(bcol);
}

```